Vuka Frontend Design

Journeys, Use Cases, Views and Wireframes

2026-09-17

1. What the frontend has to survive

This document is the frontend half of the Vuka specification. The backend spec fixes FR-1 to FR-13 and the seventeen entities. This one fixes what a person actually sees, in what order, and why.

Three rooms, three different ways to fail.

The judging room. Presentation is only 5 of the 100 marks, but every other criterion is judged through what is on the screen while we talk. Relevance, Innovation, Technical Feasibility, UX and Accessibility, Data and Insight, Security and Governance, Scalability and Sovereignty. All of them are assessed by looking. The specific risk is that someone in the room knows eQPRS, the DPME quarterly performance reporting system at eqprs.dpme.gov.za, which already configures indicators and targets, captures actuals against them, runs an approval workflow and draws dashboards. If our first screen is a capture form we have argued ourselves out of Relevance and Innovation before we open our mouths.

So the frontend leads with the thing eQPRS does not do. Every reported figure carries the cell it came from, the document it was drawn out of, and the named person who confirmed it. That is the Auditor-General's reliability test, which asks whether reported performance can be traced back to source data, rendered as an interface rather than a policy.

A three person NPO on a phone, on mobile data, in a building with one bar of signal. Six of the bodies DSAC funds are small. A reporting system they cannot use produces late submissions, which produces empty dashboards, which makes the whole thing decorative. The mobile surface is not a scaled down version of the desktop one. It is the primary surface for those users and it is built first.

A DSAC reviewer holding 28 entities and a quarter to close. Their failure mode is not confusion, it is volume. Twenty eight entities times roughly twenty targets each is a review queue no person reads end to end. The interface has to rank the queue and defend its ranking, because a reviewer who does not trust the ordering reverts to reading everything and the system has saved nobody anything.

Three rules everything else follows from

A number never appears without its provenance. Not in a tooltip, not behind a click, in the same visual unit as the number. A risk score shows its five signals. A reported actual shows its source cell. A unit cost shows its numerator and denominator. If a screen cannot show the working, the number does not go on that screen.

The interface never asks for what the system can derive. The entity is on the token, so it is never a dropdown. The quarter is derivable from the date, so it is defaulted. The deadline is derivable from the financial year, so it is displayed rather than entered. Every field we remove is a field that cannot be filled in wrong.

Every screen says what happens next and who it is waiting for. Government reporting fails at handoffs, not at data entry. A submission sitting in a state nobody owns is the normal failure. So each screen carries one line naming the current holder and the next action, and that line is generated from the workflow state rather than written into the template.

2. Who uses it

Five audiences. Four authenticate, one does not. The roles are carried as Firebase custom claims, and `entityId` on the signed token is the security boundary rather than anything the client sends.

Audience	Claim	What they want, in one sentence	What they must never be able to do
Entity reporter	`ENTITY_REPORTER` + `entityId`	Get this quarter's figures in, with proof attached, without a phone call from Pretoria in three weeks	See another entity's data. Edit a figure after confirming it. Choose which entity they are reporting for
DSAC reviewer	`DSAC_REVIEWER`	Work a queue of 28 entities in the order that matters and know which numbers they can rely on	Change a figure an entity reported. Approve a target result without seeing its source
DSAC executive	`DSAC_EXECUTIVE`	Walk into a portfolio committee meeting knowing which three entities are about to embarrass the department	Touch the data at all. This role is read only by design
Administrator	`ADMIN`	Onboard an entity, issue a reporter account, publish an entity to the citizen view	Nothing is fully barred, which is why every administrative action is written to the audit log with the actor on it
Citizen	none	Find out what a body near them was given and what it delivered	Reach anything not explicitly marked publicly visible by DSAC

The three that shape the interface

The entity reporter is the one we design for hardest, because they are the only audience whose failure breaks everything downstream. Late or evidence free submissions produce a dashboard that is empty or wrong, and an executive who stops opening it. Two sub types matter and they are genuinely different people. A finance officer at Iziko with a desktop, a template and a routine. A part time administrator at a small NPO with a phone and an hour on a Sunday. The first gets the spreadsheet upload path. The second gets the one indicator per screen mobile path. Neither is a downgrade of the other.

The DSAC reviewer is the one nobody designs for, which is why quarterly reporting in government is a folder of emailed spreadsheets. Their whole job is deciding what to look at. The queue ordering is therefore a first class feature and not a convenience, and it has to be able to answer the question “why is this one at the top” in one sentence.

The citizen view is not an afterthought and not a marketing page. It is the proof that the accountability chain closes. Allocation, target, evidence, outcome, all four visible to a person who does not work in government. It is also the only surface a judge can open on their own phone during the presentation, which is worth more than it sounds.

3. User journeys

Each journey is written as what the person is thinking, what the screen does about it, and what goes wrong if we get it lazy. The failure column is the useful one. It is where most reporting systems die.

J1. Nomsa, finance officer at a mid sized entity, desktop, quarter close

She has done this eleven times. She has the figures in her own spreadsheet already. Her fear is not the work, it is being told in October that something she submitted in July was wrong.

Step	What she does	What the screen does	Failure if we get it lazy
1	Opens Vuka, signs in	Lands on her entity’s dashboard. The entity name is already there because it is on her token. One line: “Q2 due 31 October. 14 days left. 0 of 19 targets reported”	We show a generic home page and she has to navigate to find her own entity. Every extra click is a reason to go back to email
2	Downloads the template	The template is pre filled with her 19 registered targets, their indicator codes, their annual targets and the quarter. She types actuals into one column	We hand her a blank template. She retypes indicator names, gets one wrong, and the parser cannot match it
3	Uploads the completed file	Parser reads it and shows every value next to the cell it came from, `Quarterly Report!H14` style. Nothing is saved as performance data yet	We save straight into the performance tables. Now a parse error is a data error and there is no way back
4	Reviews the parsed values	Each row has the parsed number, the source cell, a variance against the quarterly target, and a place to attach evidence. Rows with a large negative variance and no evidence are flagged before she submits, not after	We let her submit a 40 percent miss with no explanation. The reviewer sends it back and the loop costs two weeks

Step	What she does	What the screen does	Failure if we get it lazy
5	Attaches evidence	Per target. Attendance register, signed report, photograph, invoice. The interface names the Auditor-General's test in plain words: can this figure be traced back to a source document	We make evidence optional and nobody attaches any. The evidence gap signal then fires on everyone and tells us nothing
6	Confirms	A modal that says plainly: confirming writes these as the entity's reported results, in her name, and they cannot be edited afterwards. Corrections happen by the reviewer returning the submission	We make confirm a soft save. The number becomes negotiable and the audit trail becomes fiction
7	Submits	Status moves to Submitted. Screen says who holds it now and what the department does next	We show a success toast and nothing else. She does not know if anyone received it

J2. Thabo, part time administrator at a small NPO, phone, mobile data, Sunday evening

He is not a finance person. He has six indicators. He has 400MB of data left for the month and the signal drops.

Step	What he does	What the screen does	Failure if we get it lazy
1	Opens the link in an SMS	A page under 5KB. No framework, no web fonts, no images. Server rendered	We ship the React bundle to him. It costs him 2MB, takes 40 seconds, and he closes it
2	Signs in	One email field, one password field	We use a social login flow that opens a popup his browser blocks
3	Sees the list	Six indicators, each one line, each showing done or not done. Nothing else on the page	We show him the same dense table as Nomsa and he scrolls sideways forever
4	Taps the first one	One indicator, one screen. The question in plain language, the annual target, one number field, one evidence upload, one note field. Next	We put all six on one form. His connection drops on indicator five and he loses everything
5	Connection drops	Nothing is lost. The step index is in the URL, not in session state. He reopens the link and is on step 5	We hold progress in a session. The server restarts, or his session expires, and he starts again. He does not start again, he gives up

Step	What he does	What the screen does	Failure if we get it lazy
6	Finishes	A confirm screen listing all six, then done. A page that says the department has it, and when they will respond	We end on a blank page and he submits twice to be sure

J3. Lerato, DSAC reviewer, the week after the deadline

Twenty eight entities. She cannot read everything. Her real question every morning is which three things to look at today.

Step	What she does	What the screen does	Failure if we get it lazy
1	Opens the queue	Submissions ranked by risk band, not by date received. Each row carries the entity, the band, and the single highest contributing signal in words	We sort by date. She works alphabetically, runs out of week at entity 19, and the ones that mattered were at 24 and 27
2	Asks why the top one is top	Clicks the band. Panel shows all five signals, each with raw value, weight, contribution and a sentence. “Submitted an average of 34 days late across 3 prior periods”	We show a score of 71 and a red chip. She does not trust it, reverts to reading everything, and the ranking has bought nothing
3	Opens the submission	Every reported figure sits beside its source cell and its attached evidence. She can open the evidence without leaving the row	We make her download the original spreadsheet to check a number. She stops checking
4	Finds a figure she does not believe	Returns the submission with a comment on that specific target. The entity sees exactly which figure is disputed and why	We give her a single free text box for the whole submission. The entity guesses which number she means
5	Approves the rest	Approval is per submission, recorded with her name and the time	We allow silent bulk approval. The audit trail shows 28 approvals in four seconds and the Auditor-General has a question

J4. The Director-General, the day before a portfolio committee meeting

Read only. Wants three numbers and one story. Will be asked about the worst performing entity by name, in public, by someone who has read the annual report.

Step	What they do	What the screen does	Failure if we get it lazy
1	Opens the portfolio view	28 entities as a heatmap by risk band, with the totals across the top: allocated, reported, targets achieved, entities with outstanding submissions	We open on a chart nobody asked for. They want the count first
2	Looks at the worst one	Drilldown: allocation, targets, what is reported, what is not, audit history, and the risk explanation in the same words the reviewer saw	We show them a different explanation from the reviewer's. Now the department has two versions of the same fact
3	Asks what it costs	Unit cost against that entity's own plan and its own prior year. Peer comparison only within its sector, and only when a sector peer group exists	We offer cost per outcome across the portfolio. Someone asks how a ballet company compares to a boxing regulator and the credibility is gone
4	Wants it on paper	Exports a briefing note for that entity: the numbers, the signals, the evidence status and the dates	We give them a screenshot. They read from a phone in a committee meeting

J5. A citizen, no account, arriving from a link

Came from a social media post about a museum in their province. Has thirty seconds of patience.

Step	What they do	What the screen does	Failure if we get it lazy
1	Lands on the public index	Entities that DSAC has marked publicly visible, grouped by sector, each with its allocation for the year	We put a login wall in front of it, or a map that takes four seconds to draw
2	Opens one	Four things in order: what it was allocated, what it promised, what it reported, whether that was verified. In that order, because that is the accountability chain	We lead with a mission statement. They leave
3	Asks where the number came from	Each figure says which published document it came from and for which quarter	We publish unattributed numbers and become just another source to distrust

The failure mode that runs through all five

Every one of these journeys breaks at a handoff rather than at a form. Nomsa does not know if anyone received it. Thabo submits twice. Lerato's comment does not say which figure. The DG sees a different explanation from his reviewer. This is why the third design rule exists and why the state line appears on every screen.

4. Use cases

Twenty two use cases. Each one traces to a functional requirement in the backend specification, so nothing here is a screen we invented for the demo. Anything marked **cut** is the first to go if the 48 hours runs short, and Section 12 says in what order.

Entity reporter

UC-1 Download the pre filled template · FR-6 · Actor: reporter · Trigger: a reporting period opens

Main flow: reporter opens the period card, requests the template, receives an `.xlsx` carrying their entity's registered targets, indicator codes, annual targets and the period, with the actuals column empty and a hidden schema sheet naming the version. Alternate: no targets registered for the year, the system says so and names the administrator to contact rather than producing an empty file.

UC-2 Upload a completed template · FR-6 · Actor: reporter

Main flow: file is parsed, each value written as an `ExtractionResult` carrying its source cell reference. No performance data is written. Alternate A: the sheet name or header row does not match, the upload is rejected with the expected structure shown and the template offered again. Alternate B: a value is not numeric, that row is marked unparsed and the rest proceed. The reporter enters that one by hand. Alternate C: an indicator code in the file does not match a registered target, that row is held aside and shown separately rather than silently dropped.

UC-3 Confirm parsed values as reported results · FR-6, NFR-3 · Actor: reporter

Main flow: reporter reviews each extraction beside its source cell and confirms. The system writes `TargetResult` rows carrying the confirming user and the timestamp. This is the human in the loop step and it is enforced by the schema having two separate tables rather than by a policy. Alternate: reporter disagrees with a parsed value, edits it before confirming. The original extraction stays on the record beside the corrected figure.

UC-4 Attach evidence to a target · FR-1, NFR-3 · Actor: reporter

Main flow: file attached against a specific target result, stored with its uploader, its content hash and its type. The interface states the test being satisfied, which is whether the reported figure can be traced back to a source document. Alternate: evidence attached after submission. Allowed. Withholding evidence is worse than late evidence.

UC-5 Submit for review · FR-6 · Actor: reporter

Precondition: every target has either a confirmed result or an explicit reason for having none. Main flow: state moves to Submitted, the holder becomes the department, the submission timestamp is written and compared against the period due date. Alternate: submitted after the due date. Permitted, and the lateness is recorded rather than blocked. A system that refuses late submissions produces no data at all, which is the outcome we are trying to avoid.

UC-6 Respond to a returned submission · FR-7 · Actor: reporter

Main flow: reporter sees only the specific targets the reviewer disputed, with the reviewer's comment against each. Corrects and confirms again. The superseded figure and its author remain on the record.

UC-7 Report one indicator at a time on a phone · FR-10 · Actor: reporter on the mobile surface

Main flow: server rendered, one indicator per request, step index carried in the URL. Each step writes immediately. A dropped connection loses the current step and nothing else. Alternate: reporter returns days later, resumes at the first unanswered indicator.

UC-8 See the deadline and what is outstanding · FR-3 · Actor: reporter

Main flow: the entity home carries the period, the due date, days remaining and the count of targets with no result. Countdown notifications at 30 days, 15 days and in the final day.

DSAC reviewer

UC-9 Work the review queue in risk order · FR-2, FR-8 · Actor: reviewer

Main flow: submissions listed by risk band descending, each row carrying the highest contributing signal in plain language. Alternate: reviewer sorts by entity or by date. Available, and not the default.

UC-10 See why an entity scored what it scored · FR-8, NFR-4 · Actor: reviewer

Main flow: all five signals shown with raw value, normalised value, weight, contribution and a sentence. The score is never displayed without them.

UC-11 Verify a figure against its source · FR-1, FR-7 · Actor: reviewer

Main flow: each reported figure shown with its source cell and its attached evidence, openable in place. Alternate: no evidence attached. The row is marked unverifiable rather than merely unverified, which is the language the Auditor-General uses and the distinction the department cares about.

UC-12 Return a submission with per target comments · FR-7 · Actor: reviewer

Main flow: reviewer marks specific targets as disputed with a comment on each, returns the submission. Only the disputed targets are reopened.

UC-13 Approve a submission · FR-7, NFR-3 · Actor: reviewer

Main flow: approval recorded against the submission with the reviewer's identity and timestamp. Approval does not make the figures editable by anyone.

UC-14 Comment in an entity workspace · FR-4 · Actor: reviewer and reporter

Main flow: a thread per entity per period, visible to that entity and to DSAC. Replaces the email chain that currently carries this conversation and leaves no record.

DSAC executive

UC-15 See the portfolio at a glance · FR-2 · Actor: executive

Main flow: all entities by risk band with portfolio totals: allocated, reported, targets achieved against targets set, submissions outstanding.

UC-16 Drill into one entity · FR-2, FR-8 · Actor: executive

Main flow: allocation, targets, results, evidence status, audit history, risk explanation. The same explanation text the reviewer sees, from the same stored signals, because two versions of one fact is worse than none.

UC-17 Compare cost against plan and against sector peers · FR-9 · Actor: executive

Main flow: planned unit cost from allocation divided by annual target, actual from spend divided by delivery, and the variance. Peer comparison within the entity's own sector only. Explicitly absent: cost per outcome across sectors. There is no API parameter that permits it. The absence is the design decision and it is the answer to the first hard question a judge will ask about the analytics.

UC-18 Export a briefing note · FR-2 · Actor: executive · **cut candidate**

Main flow: one entity, one page, the figures and the signals and the dates, as a file that survives being read from a phone in a committee meeting.

Administrator

UC-19 Onboard an entity and issue reporter accounts · FR-5, NFR-1 · Actor: administrator

Main flow: entity created, sector assigned, reporter account issued with `role` and `entityId` set as custom claims on the Firebase token.

UC-20 Register targets for a financial year · FR-6 · Actor: administrator

Main flow: targets loaded from the entity's tabled Annual Performance Plan, each with an indicator code, an annual target and a unit. Constraint that matters: a target may only be changed where the Annual Performance Plan has been revised and re tabled, under section 4.4.4 of the 2019 Revised Framework for Strategic Plans and Annual Performance Plans. Missing a target is not a lawful reason to change it. So targets are versioned with an effective from date and a re tabling reference, and they are never mutable fields.

UC-21 Publish an entity to the citizen view · FR-11 · Actor: administrator

Main flow: `publiclyVisible` set on the entity. Nothing reaches the public surface without it. The system makes publication a switch rather than a project, and does not make the decision on the department's behalf.

Citizen

UC-22 Look up what a body received and delivered · FR-11 · Actor: anyone, no authentication

Main flow: allocation, targets, reported results, verification status, each with the published source it came from. Read only, cached, and the lightest page in the system.

5. Screen inventory

Seventeen views across three surfaces. The split between React and Thymeleaf is not a matter of taste. Anything a low bandwidth user touches is server rendered, anything an office user touches can afford a bundle.

Surface A. Low bandwidth, server rendered Thymeleaf, no JavaScript

These already exist in the repository. Every one is under 5KB before compression.

View	Route	Gate	Data	File
Mobile home	`GET /m`	reporter	open period, target count, days remaining	`mobile-home.html`
Mobile step	`GET /m/submission/{id}/step/{index}`	reporter, own entity only	one target, its annual target, prior result	`mobile-step.html`
Mobile done	after final step	reporter	submission summary and state	`mobile-done.html`
Mobile message	any dead end	reporter	one sentence and one link	`mobile-message.html`
Public index	`GET /public`	none	entities where `publiclyVisible` is true	`public-index.html`
Public entity	`GET /public/entity/{id}`	none	allocation, targets, results, sources	`public-entity.html`

Surface B. Office, React against the existing JSON API

View	Route	Gate	API it calls	State
Sign in	`/signin`	none	Firebase client SDK	build
Entity home	`/entity`	reporter	`GET /api/dashboard/entity/{entityId}`	build
Period card	`/entity/period/{id}`	reporter	`POST /api/submissions/open`	build
Template upload	`/entity/period/{id}/upload`	reporter	`POST /api/submissions/{id}/upload`	build
Extraction review	`/entity/period/{id}/review`	reporter	`GET /api/submissions/{id}/extractions`, `POST /api/submissions/{id}/confirm`	build, the most important screen in the system
Evidence drawer	overlay on extraction review	reporter	document upload endpoint	build

View	Route	Gate	API it calls	State
Submit confirm	modal	reporter	`POST /api/submissions/{id}/submit`	build
Review queue	`/review`	reviewer	`GET /api/dashboard/portfolio`	build
Submission review	`/review/{id}`	reviewer	extractions plus `POST /api/submissions/{id}/review`	build
Portfolio	`/portfolio`	reviewer, executive	`GET /api/dashboard/portfolio`	build
Entity drilldown	`/portfolio/entity/{id}`	reviewer, executive	`GET /api/dashboard/entity/{id}`, `GET /api/dashboard/entity/{id}/peers`	build
Risk panel	overlay on both	reviewer, executive	signals already returned by the dashboard endpoints	build

Surface C. Administration

View	Route	Gate	Note
Entity admin	`/admin/entities`	admin	create entity, set sector, set `publiclyVisible`
Target registration	`/admin/entities/{id}/targets`	admin	load from the tabled Annual Performance Plan, versioned, never edited in place
User admin	`/admin/users`	admin	issues the Firebase custom claims. **Cut candidate.** The claims can be set from a script for the demo

What is deliberately not a screen

There is no edit screen for a confirmed figure. Not disabled, not permission gated. It does not exist, and there is no endpoint behind it either. A reviewer who disputes a number returns the submission. This is the single easiest thing to demonstrate and the hardest thing to argue with.

There is no entity selector for a reporter. The entity comes off the signed token. A dropdown would be a way for a client to ask for another entity's data, and the absence of the control is part of the access control story rather than a UX simplification.

There is no cross sector cost comparison view. Covered in UC-17 and worth repeating here, because the absence is the point.

Route gating, exactly

Pattern	Who reaches it
`/public/**`	everyone, unauthenticated
`/m/**`	`ENTITY_REPORTER` , and the entity is taken from the token, never the path
`/api/submissions/**`	`ENTITY_REPORTER` for write, `DSAC_REVIEWER` for the review endpoint
`/api/dashboard/portfolio`	`DSAC_REVIEWER` , `DSAC_EXECUTIVE` , `ADMIN`
`/api/dashboard/entity/{id}`	the same three, plus a reporter whose token `entityId` equals `{id}`
`/admin/**`	`ADMIN`

The filter that enforces this is `FirebaseTokenFilter`, and the check that matters is the last row. A reporter reaching their own entity and nobody else's is one comparison, and it is the whole tenancy model.

6. Wireframes: the entity reporter

W1. Entity home

The entity name is present because it came off the token, not because she chose it. The state line at the bottom is generated from the workflow state and appears on every screen in the system.

```
┌──────────────────────────────────────────────────────────────────────────┐
│ VUKA         Iziko Museums of South Africa          N. Mabaso   [sign out]│
├──────────────────────────────────────────────────────────────────────────┤
│                                                                          │
│  Q2 2026/27 · 1 Jul to 30 Sep                 Due 31 Oct 2026 · 14 days left│
│                                                                          │
│  ┌────────────────────────────────────────────────────────────────────┐  │
│  │ 0 of 20 targets reported                                           │  │
│  │ ░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░░  0%     │  │
│  │                                                                    │  │
│  │  [ Download template ]   [ Upload completed file ]   [ Capture ]   │  │
│  └────────────────────────────────────────────────────────────────────┘  │
│                                                                          │
│  ▸ Waiting on you. Next: upload or capture, then confirm each figure.    │
│                                                                          │
│  Prior periods                                                           │
│  ┌────────────────────────────────────────────────────────────────────┐  │
│  │ Q1 2026/27   Approved    18 of 20 reported   submitted 2 days late │  │
│  │ Q4 2025/26   Approved    20 of 20 reported   submitted on time     │  │
│  │ Q3 2025/26   Returned    then approved       3 figures disputed    │  │
│  └────────────────────────────────────────────────────────────────────┘  │
└──────────────────────────────────────────────────────────────────────────┘
```

Why prior periods are on the home screen. The lateness signal is computed from exactly this history. Showing the reporter the same history the risk engine reads means the score is never a surprise. A system that scores you on data you cannot see is a system people work around.

W2. Extraction review

The most important screen in the product. Everything the pitch claims is visible here at once: the parsed value, the cell it came from, the variance, the evidence, and the fact that nothing has been written as performance data yet.

```
← Q2 2026/27      Reviewing parsed file: Iziko_Q2_2026.xlsx

⚠ Nothing here is saved as a reported result yet. Confirming writes these
  figures in your name, and they cannot be edited afterwards.


  HER-2.1  Number of exhibitions mounted

    Parsed    4          from  Quarterly Report!H14
    Quarter target 5     Annual target 20     Variance  −1  (−20%)

    Evidence   ✓ exhibition-schedule-q2.pdf   ✓ opening-register.pdf
               Traceable to source. [ add ]

    Note   One exhibition moved to Q3, venue works overran.
                                                        [ Confirm  4 ]


  HER-3.4  Number of learners reached by education programmes

    Parsed    2 140      from  Quarterly Report!H19
    Quarter target 3 000  Annual target 12 000  Variance  −860  (−29%)

  ⚠ No evidence attached. This figure will show to the department as
    unverifiable, which is different from unverified.       [ attach ]

    Note   ____________________________________________
    A variance past 20% needs a reason before you can confirm.
                                                    [ Confirm  2 140 ]


  HER-5.2  Heritage objects catalogued
    ✗ Not parsed. Cell Quarterly Report!H22 held "n/a".
      Enter it here, or record that there is no result this quarter.
      [ _____ ]   or   ( ) No result this quarter, because _________


  2 of 20 confirmed                           [ Confirm all remaining ]
  ▸ Waiting on you. 18 figures still to confirm before you can submit.
```

Four things this screen is doing at once. It shows provenance beside every figure, which is the reliability test. It refuses a large variance without a reason, which prevents the reviewer round trip that costs two weeks. It distinguishes unverifiable from unverified, which is the Auditor-General's own language. And it makes the irreversibility of confirming explicit before the click rather than after it.

W3. Mobile home

Under 5KB. No JavaScript. This is what a small NPO sees.

```
┌──────────────────────────┐
│ VUKA                     │
│ Luthuli Museum           │
├──────────────────────────┤
│                          │
│ Q2 2026/27               │
│ Due 31 Oct · 14 days     │
│                          │
│ 6 indicators             │
│ 0 reported               │
│                          │
│ ┌──────────────────────┐ │
│ │  Start reporting  →  │ │
│ └──────────────────────┘ │
│                          │
│ 1 Visitor numbers     ○  │
│ 2 School groups       ○  │
│ 3 Exhibitions         ○  │
│ 4 Community events    ○  │
│ 5 Objects conserved   ○  │
│ 6 Staff trained       ○  │
│                          │
│ Waiting on you.          │
└──────────────────────────┘
```

W4. Mobile step, one indicator per screen

The step index is in the URL. A dropped connection costs this screen and nothing before it.

```
┌──────────────────────────┐
│ ← Luthuli Museum   2 / 6 │
│ ▓▓▓▓▓▓░░░░░░░░░░░░░░░░░░ │
├──────────────────────────┤
│                          │
│ School groups hosted     │
│                          │
│ How many school groups   │
│ visited between 1 July   │
│ and 30 September?        │
│                          │
│ Target for the year: 40  │
│ Last quarter: 9          │
│                          │
│ ┌──────────────────────┐ │
│ │                      │ │
│ └──────────────────────┘ │
│                          │
│ Proof (optional now,     │
│ needed before approval)  │
│ [ Choose file ]          │
│                          │
│ Anything the department  │
│ should know?             │
│ ┌──────────────────────┐ │
│ │                      │ │
│ └──────────────────────┘ │
│                          │
│ ┌──────────────────────┐ │
│ │   Save and next  →   │ │
│ └──────────────────────┘ │
│                          │
│ Saved as you go. You can │
│ close this and come back.│
└──────────────────────────┘
```

The last line is doing real work. A user on a failing connection who does not know whether their work is saved will submit repeatedly or abandon. Telling them plainly, on every step, costs one line of text and removes the most common reason small entities stop using a system.

W5. Submit confirmation

```
┌──────────────────────────────────────────────────────────────┐
│  Submit Q2 2026/27 to the Department                         │
├──────────────────────────────────────────────────────────────┤
│                                                              │
│  20 targets   18 with a confirmed result                     │
│               2 recorded as having no result, with reasons   │
│               16 with evidence attached                      │
│               4 with no evidence, which the department will  │
│                 see as unverifiable                          │
│                                                              │
│  Submitting on 29 Oct 2026, two days before the due date.    │
│                                                              │
│  After this the department holds it. You will be notified if │
│  any figure is returned to you.                              │
│                                                              │
│              [ Not yet ]            [ Submit ]               │
└──────────────────────────────────────────────────────────────┘
```

7. Wireframes: the DSAC reviewer

The entities and audit outcomes in these wireframes are real and published. Robben Island Museum submitted its 2024/25 annual financial statements on 9 August 2025, missing the statutory deadline, and its audit was excluded from the portfolio outcomes as a result. Iziko and the National Library of South Africa both carry qualified opinions on heritage assets. Boxing South Africa is qualified on revenue. These are the rows the demo should show, because a judge who knows the portfolio will recognise them.

W6. Review queue

Ranked by risk, not by date received. Every row carries the one sentence that explains its position.

```
VUKA   Review queue   Q2 2026/27          L. Dlamini · DSAC   [sign out]

28 entities   19 submitted   6 outstanding   3 returned
Sorted by risk ▾     [ by entity ]  [ by date received ]

 74  CRITICAL   Robben Island Museum                    not submitted
        Largest factor: submitted an average of 41 days late across 3
        prior periods.
        Allocation R100.0m · no result on 100% of targets · audit outstanding
                                                                [ open → ]

 61  HIGH       National Library of South Africa          submitted
        Largest factor: 3 audit findings, of which 2 repeat from a prior
        year.
        Allocation R133.0m · 14 of 18 targets reported · qualified opinion
                                                                [ open → ]

 58  HIGH       Boxing South Africa                       submitted
        Largest factor: 68% of allocation drawn against 44% of target
        delivered.
        Allocation R20.7m · 16 of 18 targets reported · qualified opinion
                                                                [ open → ]

 44  MEDIUM     Iziko Museums of South Africa             submitted
        Largest factor: 4 of 20 targets have no confirmed result, with
        50% of the year elapsed.

 12  LOW        South African Heritage Resources Agency   submitted
        Largest factor: none material. Clean audit, on time, evidence
        complete.

                                          23 more · [ show all ]
```

Note what the top row is. An entity that has not submitted at all still appears in the queue, ranked, with a score. Most systems show you what arrived. The thing a reviewer needs most is what did not.

W7. Risk explanation panel

Opens from any score, anywhere in the product. Same content for the reviewer and for the Director-General, from the same stored signals.

```
Robben Island Museum                                          74 CRITICAL
Computed 17 Sep 2026 14:02 from data as at 16 Sep 2026

Submission lateness                    weight 0.30    contributes 24.6
41 days late on average across 3 prior periods.
Raw 41.0 · normalised 0.82 of a 30 day ceiling

Evidence gap                           weight 0.25    contributes 12.5
20 of 20 targets have no confirmed result, with 50% of the year
elapsed.

Spend to delivery divergence           weight 0.20    contributes 10.0
50% of allocation drawn against 0% of target delivered.

Prior audit findings                   weight 0.15    contributes 15.0
Audit outstanding. Financial statements submitted 9 Aug 2025, after
the statutory deadline. Excluded from portfolio outcomes.

Revision churn                         weight 0.10    contributes  4.0
2 targets restated mid year.

                                                   Total     74.1

Weights are fixed and published. Bands: 25 medium, 50 high, 70 critical.
This score is arithmetic, not a prediction, and can be reproduced by
hand from the figures above.
```

The last three lines are the panel’s reason for existing. A department that acts on a score has to defend it to the entity, to the portfolio committee and eventually to the Auditor-General. “Arithmetic, not a prediction” is a sentence a Director-General can say in public. No trained model produces a panel like this, which is why we did not use one.

W8. Submission review, verifying against source

```
← queue    Iziko Museums of South Africa · Q2 2026/27      44 MEDIUM ▾
Submitted 29 Oct 2026 by N. Mabaso, 2 days before the due date
16 of 20 verifiable · 4 with no evidence
Indicator            Reported  Target  Var    Source        Evidence
HER-2.1 Exhibitions       4       5    −20%   H14 ⇲       2 files  ▸
HER-3.4 Learners      2 140   3 000    −29%   H19 ⇲       none  ⚠
        ↳ confirmed by N. Mabaso, 28 Oct 2026 11:14
          ✗ Dispute this figure
          Reason (the entity sees this against this target only)
          No attendance records attached. Please attach the
          school booking register for the quarter.
HER-5.2 Objects          —       —     —     no result, reason given

1 figure disputed        [ Return with comments ]    [ Approve all ]
▸ Returning sends only the disputed target back. The rest stays as filed.
```

There is no edit control on this screen and that is the thing to point at during the demo. The reviewer can dispute, comment and return. She cannot change what the entity reported, and

neither can anyone else, because the endpoint does not exist. The ⤵ next to each source cell opens the original uploaded file at that cell.

8. Wireframes: the DSAC executive

Allocations shown are the real 2026/27 medium term estimates from Estimates of National Expenditure 2026, Vote 37, Table 37.3. The portfolio total across the 28 funded bodies is R2.258 billion, inside a vote of R6.617 billion.

W9. Portfolio

Counts first. A Director-General wants to know how many, then which ones, then why.

```
┌──────────────────────────────────────────────────────────────────────────────┐
│ VUKA    Portfolio    Q2 2026/27                Director-General · read only│
├──────────────────────────────────────────────────────────────────────────────┤
│  ┌──────────────┐ ┌──────────────┐ ┌──────────────┐ ┌──────────────┐ ┌──────────────┐  │
│  │ R2.258bn     │ │ 19 of 28     │ │ 412 of 561   │ │ 6            │ │ 3            │  │
│  │ allocated    │ │ submitted    │ │ targets      │ │ outstanding  │ │ critical     │  │
│  │ 2026/27      │ │ this quarter │ │ reported     │ │ submissions  │ │ entities     │  │
│  └──────────────┘ └──────────────┘ └──────────────┘ └──────────────┘ └──────────────┘  │
│                                                                                │
│  By risk band                        [ arts ][ heritage ][ library ][ sport ]
│  ┌──────────────────────────────────────────────────────────────────────┐  │
│  │ CRITICAL  ▌ Robben Island   ▌ Mandela Bay      ▌ Nat Heritage Council│  │
│  │           R96.8m              R35.9m              R107.7m            │  │
│  │                                                                    │  │
│  │ HIGH      ▌ Nat Library     ▌ Boxing SA        ▌ Ditsong           │  │
│  │           R137.9m            R21.5m              R124.3m            │  │
│  │           ▌ PanSALB                                                 │  │
│  │           R139.2m                                                   │  │
│  │                                                                    │  │
│  │ MEDIUM    ▒ Iziko  ▒ PACOFS  ▒ Freedom Park  ▒ State Theatre       │  │
│  │           ▒ NFVF   ▒ Playhouse  ▒ Nelson Mandela Museum  + 4       │  │
│  │                                                                    │  │
│  │ LOW       ░ SAHRA  ░ SAIDS  ░ Market Theatre  ░ Artscape  ░ NAC    │  │
│  │           ░ Luthuli  ░ Amazwi  ░ War Museum  ░ WHAG  + 2           │  │
│  └──────────────────────────────────────────────────────────────────────┘  │
│                                                                                │
│  ▸ R358.6m sits with entities in the critical band. Click any name for the │
│    figures behind its score.                                                │
└──────────────────────────────────────────────────────────────────────────────┘
```

The bottom line is the one that gets quoted. Not “three entities are critical” but “R358.6 million sits with entities in the critical band”. Risk expressed in rands is the sentence a Director-General can take into a committee meeting, and it is one multiplication away from data we already hold.

W10. Entity drilldown

```
┌──────────────────────────────────────────────────────────────────────────┐
│ ← portfolio    National Library of South Africa          █ 61 HIGH  ▾    │
│                 Schedule 3A · library sector · reports to the Minister   │
├──────────────────────────────────────────────────────────────────────────┤
│ THE CHAIN                                                                │
│                                                                          │
│  Allocated          Promised          Reported         Verified          │
│  R137 898 000  →   18 targets   →    14 of 18    →   9 of 14            │
│  ENE 2026            APP 2026/27       Q1 and Q2       evidence attached │
│  Table 37.3          tabled 2026-03    submitted                         │
│                                                                          │
│  ──────────────────────────────────────────────────────────────────────  │
│ TRAJECTORY                                Allocation, R million          │
│    2023/24  ██████████████████████████  128.0                           │
│    2024/25  ██████████████████████████  128.8                           │
│    2025/26  ███████████████████████████  133.0                          │
│    2026/27  ████████████████████████████  137.9                         │
│                                                                          │
│ AUDIT HISTORY                                                            │
│    2024/25   Qualified   heritage assets not properly accounted for      │
│    2023/24   Qualified   same basis, repeat finding                      │
│    AGSA: "cycle of qualified audit outcomes"                             │
│                                                                          │
│ OUTSTANDING                                                              │
│    4 targets with no result at half year                                 │
│    5 reported figures with no evidence attached                          │
│                                                                          │
│    [ Explain the score ]   [ Unit cost ]   [ Briefing note ]             │
└──────────────────────────────────────────────────────────────────────────┘
```

The four boxes at the top are the whole product in one row. Allocated, promised, reported, verified. Each one cites where it came from. That row is what the pitch opens on, and it is the reason the accountability chain is modelled as foreign keys rather than as a diagram on a slide.

W11. Unit cost

```
┌──────────────────────────────────────────────────────────────────────────┐
│ Unit cost   National Library of South Africa                             │
│ Indicator LIB-1.2  Items digitised and made publicly accessible          │
├──────────────────────────────────────────────────────────────────────────┤
│                                                                          │
│ AGAINST ITS OWN PLAN                                                     │
│   Planned    R24 000 000 apportioned ÷ 12 000 items      =  R2 000       │
│   Actual     R13 100 000 spent      ÷  4 900 delivered   =  R2 673       │
│   Variance   +33.7% per item                                             │
│                                                                          │
│ AGAINST ITS OWN HISTORY                                                  │
│   2024/25    R2 410 per item                                             │
│   2023/24    R2 190 per item                                             │
│                                                                          │
│ AGAINST SECTOR PEERS        library sector, 2 comparable entities        │
│   This entity   R2 673                                                   │
│   Peer median   R2 240                                                   │
│                                                                          │
│  ┌──────────────────────────────────────────────────────────────────┐   │
│  │ Comparison is within the library sector only. There is no view that │ │
│  │ compares cost per outcome across sectors, because a library item and │ │
│  │ a boxing licence are not commensurable outputs. The system does not │ │
│  │ offer that comparison and there is no parameter that produces it.   │ │
│  └──────────────────────────────────────────────────────────────────┘   │
└──────────────────────────────────────────────────────────────────────────┘
```

Putting the limitation on the screen rather than in a footnote is what makes the analytics credible. Every judge has seen a dashboard that compares incomparable things. Showing that we thought about it, and built the restriction into the API rather than into a guideline, converts the weakest part of any performance analytics story into the strongest.

The figures inside this wireframe are illustrative apportionments. Section 12 says when they get replaced.

9. Wireframes: the citizen view

The only surface with no login, and the only one a judge can open on their own phone while we are still talking. It is also the smallest page in the system, which is not a coincidence: it has to load on a bad connection for someone with no reason to be patient.

Nothing appears here unless DSAC has set `publiclyVisible` on that entity. Every seeded entity ships with it false, so the department turns this on rather than discovering it is already on.

W12. Public index

```
┌──────────────────────────────────────────────────────────┐
│  VUKA                                                    │
│  What the bodies funded by Sport, Arts and Culture       │
│  received, and what they delivered.                      │
├──────────────────────────────────────────────────────────┤
│                                                          │
│  2026/27 · figures from Estimates of National Expenditure │
│                                                          │
│  HERITAGE                                                │
│  ────────────────────────────────────────────────────  │
│  Ditsong Museums of South Africa              R124 346 000  ›  │
│  Freedom Park                                 R114 075 000  ›  │
│  Iziko Museums of South Africa                R110 607 000  ›  │
│  Robben Island Museum                          R96 799 000  ›  │
│  South African Heritage Resources Agency       R74 953 000  ›  │
│  Luthuli Museum                                R19 369 000  ›  │
│                                                          │
│  LIBRARIES                                               │
│  ────────────────────────────────────────────────────  │
│  National Library of South Africa             R137 898 000  ›  │
│  South African Library for the Blind           R39 220 000  ›  │
│                                                          │
│  ARTS                                                    │
│  ────────────────────────────────────────────────────  │
│  National Film and Video Foundation           R338 047 000  ›  │
│  National Arts Council                        R297 212 000  ›  │
│  Market Theatre Foundation                     R57 445 000  ›  │
│                                                          │
│  ────────────────────────────────────────────────────  │
│  Only bodies the Department has chosen to publish appear here. │
│  Allocations are from Vote 37, Table 37.3.               │
└──────────────────────────────────────────────────────────┘
```

The last two lines matter more than they look. The first tells a citizen that this is not the full portfolio, which stops the page from being read as a complete picture when it is not. The second cites the source, which is what separates this from an unattributed number on a government website that nobody believes.

W13. Public entity page

The accountability chain, in the order a person who does not work in government would ask about it.

```
‹ all bodies

Luthuli Museum
Heritage · Stanger, KwaZulu-Natal

WHAT IT WAS GIVEN
R19 369 000 for 2026/27
Estimates of National Expenditure 2026, Vote 37, Table 37.3

  2023/24   R17 689 000
  2024/25   R17 814 000
  2025/26   R18 601 000
  2026/27   R19 369 000

WHAT IT PROMISED
6 targets for the year, from its Annual Performance Plan

WHAT IT REPORTED
Up to 30 September 2026

  Visitors received            4 210  of 8 000 for the year
  School groups hosted            17  of 40
  Exhibitions mounted              2  of 4
  Community events                 3  of 6
  Objects conserved               no result reported yet
  Staff trained                    4  of 8

WAS IT CHECKED
5 of 6 figures have supporting evidence on file and were
confirmed by a named official at the museum.
1 figure has not been reported.

This is not an audit. The Auditor-General audits these
bodies annually and those reports are tabled in Parliament.

Last updated 31 October 2026, after the Department reviewed
the second quarter submission.
```

Three things this page deliberately does not do

No risk score. The score is an internal management tool with weights we chose. Publishing it would turn a departmental prioritisation aid into a public judgement about a named institution, which is not a thing we get to do on DSAC's behalf and not a thing the arithmetic was built to carry.

No league table. Entities are listed by sector and by size of allocation, not ranked. The moment this page ranks, it will be screenshotted as a ranking, and a museum and a film fund will be compared on numbers that do not compare.

No claim to be an audit. The line saying so is not legal cover, it is accuracy. Confirmed by a named official with a document attached is a real and useful standard, and it is a different standard from an audit opinion. Blurring the two would be the fastest way to lose the Auditor-General's people, who will be in the room.

The "no result reported yet" line stays visible. A public page that quietly omits what was not done is worse than no public page. That single line is the reason this surface belongs in the product rather than in the marketing.

10. Components and states

The screens above reuse nine components. Building these nine first and assembling screens from them is what makes a two day frontend possible, and it is also why the reviewer and the Director-General see the same risk explanation rather than two versions of one fact.

Design tokens

Deliberately small. A palette this tight is hard to make ugly under time pressure.

```
ink             #1A1A1A    body text
ink-muted       #5A6B64    secondary text, citations
rule            #D8E0DC    borders, dividers
canvas          #FFFFFF    page
canvas-sunk     #F4F8F6    cards, table stripes
brand           #0F3D2E    headings, primary buttons
brand-light     #17614A    links, secondary headings

band-low        #2E7D5B    risk bands, and the only place
band-medium     #B8860B    colour carries meaning
band-high       #C1571A
band-critical   #A32E2E

type            system sans, 16px base, 1.55 line height
space           4 8 12 16 24 32 48
radius          4px, everywhere, no exceptions
```

Risk band colour is never the only carrier. Every band shows its numeric score and its label in text, because roughly eight percent of men have some colour vision deficiency and because a printed briefing note goes out in greyscale. Where a band appears as a block character in these wireframes, the shipped version uses a filled swatch plus the number plus the word.

The nine components

Component	Where it appears	The rule it carries
`StateLine`	every screen, bottom	Names the current holder and the next action. Generated from workflow state, never written into a template
`RiskBadge`	queue, portfolio, drilldown	Score, band word, band colour. Opens `RiskPanel` on click. Never appears alone
`RiskPanel`	overlay, three screens	Five signals with raw value, weight, contribution, sentence. Reads stored signals, never recomputes
`ProvenanceCell`	extraction review, submission review, export	A figure plus its source cell plus a link that opens the original file at that cell
`EvidenceChip`	wherever a figure appears	File count, AGSA criterion satisfied, or the unverifiable warning. Three states, one component
`ChainStrip`	entity drilldown, public entity page	Allocated, promised, reported, verified, each with its citation
`PeriodCard`	entity home, mobile home	Period, due date, days remaining, deadline basis in plain words

Component	Where it appears	The rule it carries
`IndicatorRow`	every list of targets	Indicator, target, actual, variance, status. The unit of work across the whole product
`CitationLine`	under any published figure	Document, table, date. Small, grey, always present

States, which is where frontends actually fail

Every component ships with five states. A demo that only has the happy path breaks live, and the failure is always the same shape: an empty database, a slow request, or a permission the presenter forgot to set.

Component	Empty	Loading	Error	Denied
`RiskBadge`	“Not yet scored” in grey, no colour	skeleton pill	“Score unavailable” and the panel still opens on stored signals	not rendered
`RiskPanel`	“No signals stored for this period”	skeleton rows	names which signal failed to load, shows the rest	not rendered
`ProvenanceCell`	“Entered by hand, no source cell”	value first, cell reference second	value shown, link disabled with a reason	value hidden entirely
`EvidenceChip`	“No evidence attached”, warning styling	spinner on the count	“Could not read the document store”	count shown, file not openable
`ChainStrip`	each box independently shows a dash and its reason	per box	per box, never the whole strip	whole strip hidden
`IndicatorRow`	“No result reported” rather than a blank or a zero	skeleton	value preserved, status marked unknown	row absent
`PeriodCard`	“No open reporting period” plus who to contact	skeleton	“Could not load the period” plus a retry	not rendered

Empty is never zero. A target with no result reported shows as “no result reported”, not as 0. The difference between not done and not reported is the entire subject of this product, and a component that renders both as `0` destroys it at the last inch.

Denied is never an error. A reporter who somehow reaches another entity’s submission gets a not-found, not a forbidden. Telling a caller that a record exists but is not theirs is itself a disclosure about another entity, and the backend already behaves this way in `ExportController`. The frontend has to match it rather than helpfully explaining.

One interaction rule

Nothing irreversible happens without a sentence saying so, in the same click. Confirming a figure and submitting a period are the two irreversible actions in the product. Both carry a plain statement of what becomes permanent and whose name goes on it, in the modal rather than in a tooltip, before the button rather than after it. Everything else can be a single click with no ceremony, which is what buys the attention for these two.

11. Accessibility and bandwidth

UX and Accessibility is 10 of the 100 marks, and it is the criterion most teams treat as a checkbox. It is also the one where a specific, measured claim beats a general one by a wide margin.

The page weight budget

Measured, not aspirational. The six Thymeleaf templates in the repository are already under budget and their real sizes are below.

Surface	Budget	Actual	What is in it
`mobile-message.html`	5KB	**617 B**	one sentence, one link
`mobile-done.html`	5KB	**1 873 B**	submission summary
`public-index.html`	5KB	**2 187 B**	every published entity, grouped
`mobile-home.html`	5KB	**2 312 B**	period, six indicators, progress
`public-entity.html`	5KB	**3 918 B**	the full accountability chain
`mobile-step.html`	5KB	**4 915 B**	one indicator, input, upload, note
React dashboard	250KB gzipped	to build	office users only, never on the reporter path

No web fonts, no icon fonts, no framework, no images on any low-bandwidth surface. Inline CSS, because a separate stylesheet is a second request and on a bad connection a second request is the one that fails.

What that buys, stated the way it should be said out loud. On a 2G connection at roughly 50 kbit/s, `mobile-step.html` arrives in under a second. A 2MB React bundle takes about five minutes and will usually fail first. Six of the bodies DSAC funds are small enough that this is the difference between reporting and not reporting.

Accessibility, specifically

WCAG 2.1 AA, and the parts that actually get tested:

Requirement	How it is met
Contrast 4.5:1 on body text	`ink` on `canvas` is 15.9:1. `ink-muted` on `canvas` is 5.4:1. Both pass
Contrast 3:1 on large text and UI borders	`brand` on `canvas` is 11.2:1
Colour is not the only carrier	every risk band shows score, word and swatch together
Keyboard reachable	server-rendered surfaces are plain forms and links, so this is free rather than engineered
Focus visible	2px `brand-light` outline, never removed
Labels tied to inputs	explicit `for` and `id` on every field, no placeholder-as-label anywhere
Errors announced	error text sits next to its field and is referenced by `aria-describedby`, not shown only as a red border
Language declared	`lang="en-ZA"` on every page
Zoom to 200%	single column layouts, no fixed heights, no horizontal scroll
Touch targets	44px minimum on every mobile control

Two claims we do not make

We do not claim multilingual support. South Africa has twelve official languages and a performance reporting system that ships in English only is a real limitation. The honest position is that the interface strings are externalised so translation is a content task rather than a rebuild, and that PanSALB, which is in the seeded portfolio, would be the obvious partner. Claiming twelve languages in a 48 hour build and then demonstrating English would be worse than saying this.

We do not claim screen reader testing. The markup is built to be screen reader friendly and nobody has run it through NVDA or VoiceOver. Say built for, not tested with. An accessibility claim that fails a live check is worse than a smaller claim that holds.

Why the mobile surface carries the inclusivity argument alone

WhatsApp was the original plan for reaching small entities and it is not achievable here. Meta’s business verification takes weeks and cannot be compressed for a hackathon. Rather than demo a mock WhatsApp flow and hope nobody asks, the low-bandwidth web surface ships as the real inclusivity answer, it is measured, and it works on the judge’s own phone during the presentation. `SubmissionChannel` retains a `WHATSAPP` value so the path stays open, and that is the extent of the claim.

12. Build order, and what gets cut

The six Thymeleaf surfaces already exist. Everything below is the React dashboard and the reporter’s desktop path.

Order

The sequence is chosen so that stopping at the end of any block still leaves something demonstrable. That is the only property that matters when the clock is real.

Block	Hours	What gets built	Demonstrable at the end
1	3	Nine components with all five states, against stubbed data	A component page. Nothing user facing, and skipping it costs more than it saves
2	3	Sign in, route gating, `VukaPrincipal` wired through	Four roles land in four different places, and a reporter cannot reach the portfolio
3	4	**Extraction review**, the screen the pitch turns on	Upload a template, see every figure beside its source cell, confirm, submit
4	3	Review queue and submission review	A reviewer works the queue in risk order and returns one disputed figure
5	3	Portfolio and entity drilldown	The Director-General's view, with `ChainStrip` carrying real allocations
6	2	`RiskPanel` on all three screens	Any score anywhere opens the same explanation
7	2	Unit cost view	Including the absent comparison, stated on screen
8	2	Polish, empty states verified, 200% zoom pass, dry run	The demo, twice, end to end

Twenty two hours of frontend inside the 48. The remainder is backend integration, the seed data pass and the presentation itself.

Block 3 is the one to protect. If everything after it is rough, the product still demonstrates its central claim, which is that a reported figure carries its provenance and a human confirmed it. If block 3 is rough and everything else is polished, we have built a nicer version of something DPME already runs.

Cut order

If the clock beats us, cut from the bottom of this list up. Decided now, in the cold, rather than at 3am.

Cut	Cost of cutting	Why it goes first
1. User admin screen	none for the demo	Firebase claims can be set from a script. Nobody watches an admin form

Cut	Cost of cutting	Why it goes first
2. Briefing note export	low	The CSV export already exists in the backend and covers the requirement
3. Target registration screen	low	Targets are seeded. This is a data loading task, not a story
4. Unit cost view	**medium, and be careful here**	The absent cross-sector comparison is a strong answer to a question judges ask. If cut, the point has to be made verbally and in the specification
5. Entity workspace comments	medium	Real requirement, weak demo. The threading is visible in the schema
6. Portfolio sector filters	low	Twenty eight rows fit on one screen

Never cut: the `RiskPanel`, the provenance cell, the mobile flow, or the public entity page. Those four are the four judging criteria we score best on, and each one is the visible form of a decision made in the backend rather than a screen for its own sake.

The demo path, in order

Eight minutes, and it follows the accountability chain rather than the menu structure.

1. **Open the public entity page on a phone.** No login. Allocation, promise, report, verification, with citations. Thirty seconds, and it establishes that the chain closes before we explain anything.
2. **Robben Island in the review queue, at the top, not submitted.** Open the risk panel. Read the sentence about the statutory deadline aloud. This is a real entity, a real missed date and a real exclusion from the portfolio audit outcomes.
3. **Upload a template as a reporter.** Show the parsed figures beside their source cells. Point out that nothing is saved as performance data yet.
4. **Confirm one figure.** Show the modal saying it cannot be edited afterwards. Then try to find an edit control anywhere in the reviewer's view, and fail, because there is no endpoint behind one.
5. **Submit on a phone.** One indicator per screen, under 5KB. Hand the phone to a judge.
6. **Portfolio view.** "R358.6 million sits with entities in the critical band."
7. **Unit cost.** Show the box explaining the comparison we refuse to offer.
8. **Export.** The full CSV with provenance columns, then the eQPRS shape without them. Close on the line that makes the positioning land: eQPRS is where the numbers go, this is where they can be defended.

What to say when asked the hard question

Someone will ask how this differs from eQPRS. The answer is short and it is in the README: eQPRS captures, it does not hold evidence, and it was built for departments while every DSAC body is Schedule 3A, reporting to its executive authority under a regulation that sets no deadline and names no system. We fill that gap and we export to theirs.

And when asked what is not real: the quarterly submission timings in the seed are illustrative because that data is not published at this granularity. The allocations, the audit outcomes and the targets achieved are real and cited, and the risk engine reproduces the Auditor-General’s own conclusions from them. Saying this before being asked is worth more than being caught not saying it.